

CS 273A Machine Learning (Fall 2017) Homework 1

October 9, 2017

Author: Chenxi Wang

```
In [229]: import numpy as np
          import matplotlib.pyplot as plt
```

```
In [8]: import os
        os.getcwd()
```

```
Out[8]: '/Users/sheilacwang/Documents/Study /17Fall/CS 273A/hw/hw1'
```

1 Problem 1: Python & Data Exploration

```
In [217]: iris = np.genfromtxt("data/iris.txt", delimiter=None)
```

```
In [371]: Y = iris[:, -1] # target value is the last column
          X = iris[:, 0:-1] # features are the other columns
```

1.1 Question 1

```
In [221]: print X.shape
```


```
(148, 4)
```

There are 148 data points and each one has 4 features.

1.2 Question 2

```
In [349]: fig, ax = plt.subplots(2, 2, figsize=(10, 8))
          ax = ax.ravel()
          axes = ['Sepal Length', 'Sepal Width', 'Petal Length', 'Petal Width']

          for j, ax in enumerate(ax):
              ax.hist(X[:, j])
              ax.set_title('Histogram of Feature %s: %s'%(j+1, axes[j]))
              ax.set_xlabel(axes[j])
              ax.set_ylabel('Frequency')
          plt.tight_layout()
          plt.show()
```



1.3 Question 3

```
In [23]: np.set_printoptions(precision=4)
```

```
In [306]: print 'The mean of the data points for each feature:'
          print np.mean(X, axis=0)
```

The mean of the data points for each feature:
[5.9001 3.0989 3.8196 1.2526]

```
In [307]: print 'The standard deviation of the data points for each feature:'
          print np.std(X, axis=0)
```

The standard deviation of the data points for each feature:
[0.8334 0.4363 1.7541 0.7588]


1.4 Question 4

```
In [381]: fig, ax = plt.subplots(1,3, figsize=(12, 4))
          ax = ax.ravel()
          xaxes = ['Sepal Length', 'Sepal Width', 'Petal Length', 'Petal Width']
          crs = ['green', 'yellow', 'orange']

          for j,ax in enumerate(ax):

              for c in np.unique(Y):
                  ax.plot(X[Y==c, 0], X[Y==c, j+1], 'o', color = crs[int(c)], \
                          label='class %d' % (int(c)+1))

              ax.set_title('Scatterplot of Feature 1 vs %s' %(j+2))
              ax.set_xlabel(xaxes[0])
              ax.set_ylabel(xaxes[j+1])
              ax.legend()
          plt.tight_layout()
          plt.show()
```



2 Problem 2: KNN Predictions

```
In [389]: import mltools as ml

          Y = iris[:, -1] # target value is the last column
          X = iris[:, 0:2] # features are the first 2 columns

In [390]: np.random.seed(0)

          X,Y = ml.shuffleData(X,Y); # shuffle data randomly
          Xtr,Xva,Ytr,Yva = ml.splitData(X,Y, 0.75);
          # split data into 75/25 train/validation
```

2.1 Question 1


The visualization plots of the training data classification are given below:

```
In [409]: K = [1, 5, 10, 50];


for i, k in enumerate(K):
    knn = ml.knn.knnClassify()
    knn.train(Xtr, Ytr, k)
    YvaHat = knn.predict(Xva)
    fig, ax = plt.subplots(1, 1, figsize=(4, 3))
    ml.plotClassify2D(knn, Xtr, Ytr)
    ax.set_title(
        'KNN Classification on Training Data with K = %s' %k
    )

plt.show()
```


KNN Classification on Training Data with K = 1




KNN Classification on Training Data with $K = 5$



KNN Classification on Training Data with $K = 10$



KNN Classification on Training Data with K = 50




Next I'll visualize the classification boundaries using the validation data:


```
In [410]: K = [1, 5, 10, 50];
```

```
for i, k in enumerate(K):
    knn = ml.knn.knnClassify()
    knn.train(Xtr, Ytr, k)
    YvaHat = knn.predict(Xva)
    fig, ax = plt.subplots(1, 1, figsize=(4, 3))
    ml.plotClassify2D( knn, Xva, Yva )
    ax.set_title(
        'KNN Classification on Validation Data with K = %s' %k
    )
    plt.show()
```


KNN Classification on Validation Data with K = 1



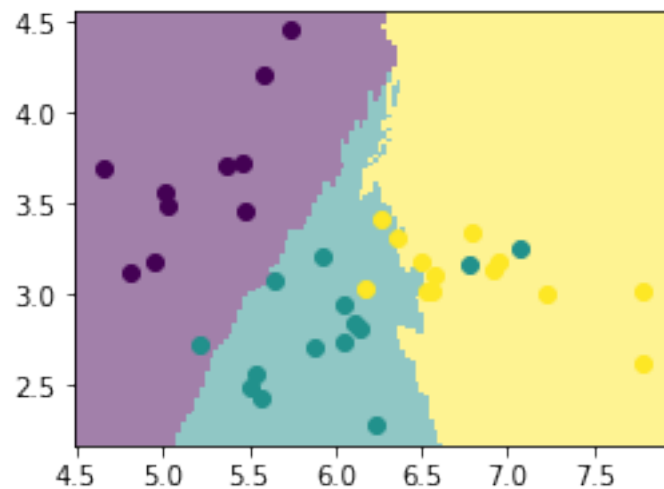
KNN Classification on Validation Data with $K = 5$



KNN Classification on Validation Data with $K = 10$



KNN Classification on Validation Data with K = 50



2.2 Question 2

Using only the first two features, error rates on both the training data and validation data for each value of k are computed, and a semi-log plot is shown as following:

```
In [351]: Y = iris[:, -1] # target value is the last column
          X = iris[:, 0:2] # features are the first 2 columns

          np.random.seed(0)
          X,Y = ml.shuffleData(X,Y); # shuffle data randomly
          Xtr,Xva,Ytr,Yva = ml.splitData(X,Y, 0.75);

          K = [1,2,5,10,50,100,200];
          errTrain = [];
          errVa = [];

          fig, ax = plt.subplots(1, 1, figsize=(5, 4))

          for i,k in enumerate(K):
              learner = ml.knn.knnClassify()
              k = K[i]
              print 'k =', k
              learner.train(Xtr, Ytr, k)
              Yhat = learner.predict(Xtr)
              errTrain.append(np.mean(Yhat != Ytr))
              YvaHat = learner.predict(Xva)
              errVa.append(np.mean(YvaHat != Yva))
```



```
print "Training Error Rate:", errTrain[i]
print "Validation Error Rate:", errVa[i]
print
```

```
ax.semilogx(K, errTrain, "r", label = "Training Error Rate")
ax.semilogx(K, errVa, "g", label = "Validation Error Rate" )
ax.legend()
plt.show()
```

```
k = 1
Training Error Rate: 0.0
Validation Error Rate: 0.297297297297
```

```
k = 2
Training Error Rate: 0.117117117117
Validation Error Rate: 0.297297297297
```

```
k = 5
Training Error Rate: 0.135135135135
Validation Error Rate: 0.27027027027
```

```
k = 10
Training Error Rate: 0.18018018018
Validation Error Rate: 0.378378378378
```

```
k = 50
Training Error Rate: 0.261261261261
Validation Error Rate: 0.135135135135
```

```
k = 100
Training Error Rate: 0.414414414414
Validation Error Rate: 0.324324324324
```

```
k = 200
Training Error Rate: 0.648648648649
Validation Error Rate: 0.72972972973
```



Based on these results, the value of k I recommend is 50 because it has the least error rate on the validation data.

2.3 Question 3

In [352]: *## With all the features:*

```
Y = iris[:, -1] # target value is the last column
X = iris[:, 0:-1] # features are the first 4 columns

np.random.seed(0)
X, Y = ml.shuffleData(X, Y); # shuffle data randomly
Xtr, Xva, Ytr, Yva = ml.splitData(X, Y, 0.75);

K = [1, 2, 5, 10, 50, 100, 200];
errTrain = [];
errVa = [];

fig, ax = plt.subplots(1, 1, figsize=(5, 4))

for i, k in enumerate(K):
    learner = ml.knn.knnClassify()
    k = K[i]
    print 'k =', k
    learner.train(Xtr, Ytr, k)
```

```

        Yhat = learner.predict(Xtr)
        errTrain.append(np.mean(Yhat != Ytr))
        YvaHat = learner.predict(Xva)
        errVa.append(np.mean(YvaHat != Yva))
        print "Training Error Rate:", errTrain[i]
        print "Validation Error Rate:", errVa[i]
        print

    ax.semilogx(K, errTrain, "r", label = "Training Error Rate")
    ax.semilogx(K, errVa, "g", label = "Validation Error Rate" )
    ax.legend()
    plt.show()

k = 1
Training Error Rate: 0.0
Validation Error Rate: 0.0540540540541

k = 2
Training Error Rate: 0.027027027027
Validation Error Rate: 0.027027027027

k = 5
Training Error Rate: 0.018018018018
Validation Error Rate: 0.027027027027

k = 10
Training Error Rate: 0.018018018018
Validation Error Rate: 0.027027027027

k = 50
Training Error Rate: 0.117117117117
Validation Error Rate: 0.0540540540541

k = 100
Training Error Rate: 0.378378378378
Validation Error Rate: 0.378378378378

k = 200
Training Error Rate: 0.648648648649
Validation Error Rate: 0.72972972973

```



Now with all the features in the data set, the semi-log plot is very different from the previous one, and this time the k value of 5 and 10 brings the least error rate on both training data and validation data.

3 Problem 3: Naive Bayes Classifiers

3.1 Question 1

The class probability:

$$P(y = 1) = \frac{4}{10} = \frac{2}{5}; P(y = -1) = \frac{6}{10} = \frac{3}{5}$$

The individual feature probabilities for each class y and feature X_i :

$$P(X_1 = 0|y = 1) = \frac{1}{4}; P(X_1 = 1|y = 1) = \frac{3}{4}$$

$$P(X_1 = 0|y = -1) = \frac{1}{2}; P(X_1 = 1|y = -1) = \frac{1}{2}$$

$$P(X_2 = 0|y = 1) = 1; P(X_2 = 1|y = 1) = 0$$

$$P(X_2 = 0|y = -1) = \frac{1}{6}; P(X_2 = 1|y = -1) = \frac{5}{6}$$

$$P(X_3 = 0|y = 1) = \frac{1}{4}; P(X_3 = 1|y = 1) = \frac{3}{4}$$

$$P(X_3 = 0|y = -1) = \frac{1}{3}; P(X_3 = 1|y = -1) = \frac{2}{3}$$

$$\begin{aligned}
P(X_4 = 0|y = 1) &= \frac{1}{2}; P(X_4 = 1|y = 1) = \frac{1}{2} \\
P(X_4 = 0|y = -1) &= \frac{1}{6}; P(X_4 = 1|y = -1) = \frac{5}{6} \\
P(X_5 = 0|y = 1) &= \frac{3}{4}; P(X_5 = 1|y = 1) = \frac{1}{4} \\
P(X_5 = 0|y = -1) &= \frac{2}{3}; P(X_5 = 1|y = -1) = \frac{1}{3}
\end{aligned}$$

3.2 Question 2

For $X = [0, 0, 0, 0, 0]$, the predicted class is $y = +1$; For $X = [1, 1, 0, 1, 0]$, the predicted class is $y = -1$.

3.3 Question 3

$$P(y = 1|X = [1, 1, 0, 1, 0]) = \frac{P(X = [1, 1, 0, 1, 0]|y = 1)P(y = 1)}{P(X = [1, 1, 0, 1, 0]|y = 1)P(y = 1) + P(X = [1, 1, 0, 1, 0]|y = -1)P(y = -1)}$$

Since $P(X_2 = 1|y = 1) = 0$, therefore the posterior probability is 0.

3.4 Question 4

Because our data samples are less than the number of parameters 2^5 . For some feature combinations that we don't observe in the training data, the "joint" Bayes classifier would predict 0 probability in the validation data, which could result in misclassification in this case since for ties we prefer to predict class +1. And using more complex models would cause overfitting.

3.5 Question 5

The naive Bayes model over features x_2, x_3, x_4, x_5 is:

$$P(y = 1|X = [x_2, x_3, x_4, x_5]) =$$

$$\frac{P(X = [x_2, x_3, x_4, x_5]|y = 1)P(y = 1)}{P(X = [x_2, x_3, x_4, x_5]|y = 1)P(y = 1) + P(X = [x_2, x_3, x_4, x_5]|y = -1)P(y = -1)},$$

where

$$P(X = [x_2, x_3, x_4, x_5]|y = 1) = P(x_2|y = 1)P(x_3|y = 1)P(x_4|y = 1)P(x_5|y = 1)$$

4 Statement of Collaboration

I have abided by the rules of conduct and academic honesty adopted by UC Irvine. I did not discuss the specific solutions to this homework with any person.

Chenxi Wang 10/8/2017